





```
In [ ]: pred = dt.predict(X_test)
print(classification_report(y_test, pred))

              precision    recall  f1-score   support

    0         0.92         0.90         0.91         1464
    1         0.80         0.84         0.82          723

 accuracy         0.86         0.87         0.88         2187
 macro avg         0.86         0.87         0.86         2187
 weighted avg         0.88         0.88         0.88         2187
```

## 6.XGBoost

```
In [ ]: xgb_cl = xgb.XGBClassifier()

xgb_cl.fit(X_train_new, y_train_new)
print("Accuracy on Train data: ",xgb_cl.score(X_train_new, y_train_new))

Accuracy on Train data: 0.981957353745216
```

```
In [ ]: pred = xgb_cl.predict(X_val)
print(classification_report(y_val, pred))

              precision    recall  f1-score   support

    0         0.96         0.94         0.95         2195
    1         0.88         0.91         0.89         1084

 accuracy         0.92         0.92         0.93         3279
 macro avg         0.92         0.92         0.92         3279
 weighted avg         0.93         0.93         0.93         3279
```

```
In [ ]: pred = xgb_cl.predict(X_test)
print(classification_report(y_test, pred))

              precision    recall  f1-score   support

    0         0.96         0.94         0.95         1464
    1         0.89         0.91         0.90          723

 accuracy         0.93         0.93         0.93         2187
 macro avg         0.92         0.93         0.93         2187
 weighted avg         0.93         0.93         0.93         2187
```

## 7.Random Forest Tree

```
In [ ]: rf = RandomForestClassifier(n_estimators = 100, random_state=0)
rf.fit(X_train_new, y_train_new)
print("Accuracy on Train data: ",rf.score(X_train_new, y_train_new))

Accuracy on Train data: 0.998587570621469
```

```
In [ ]: pred = rf.predict(X_val)
print(classification_report(y_val, pred))

              precision    recall  f1-score   support

    0         0.94         0.94         0.94         2195
    1         0.87         0.89         0.88         1084

 accuracy         0.91         0.91         0.91         3279
 macro avg         0.91         0.91         0.91         3279
 weighted avg         0.92         0.92         0.92         3279
```

```
In [ ]: pred = rf.predict(X_test)
print(classification_report(y_test, pred))

              precision    recall  f1-score   support

    0         0.94         0.94         0.94         1464
    1         0.88         0.88         0.88          723

 accuracy         0.92         0.92         0.92         2187
 macro avg         0.91         0.91         0.91         2187
 weighted avg         0.92         0.92         0.92         2187
```

## Final Model Building(Hyperparamter Tuning)

From the results above we saw that the XGBoost model gave us the highest f1 and recall, we will now do some paramter tuning and then make final conclusions which parameters should be used based on the results for future detection.

### XGBoost

```
In [ ]: parama={
    "learning_rate": [0.05, 0.10, 0.20, 0.30 ],
    "max_depth": [ 3, 5, 8, 10, 15],
    "min_child_weight": [ 1, 3, 5, 7 ],
    "gamma": [ ( 0.0, 0.1, 0.2 , 0.3, 0.4 ),
    "colsample_bytree": [ 0.5, 0.4, 0.5 , 0.7 ]
}

classifier = XGBClassifier()
grid_search = GridSearchCV(estimator = classifier,param_grid=params,scoring='f1',n_jobs=-1,cv=3, verbose=1)
grid_search.fit(X_train_new,y_train_new.ravel())
grid_search.best_params_

Fitting 3 folds for each of 3200 candidates, totalling 9600 fits
('colsample_bytree': 0.5,
 'gamma': 0.2,
 'learning_rate': 0.2,
 'max_depth': 15,
 'min_child_weight': 1,
 'n_estimators': 100,
 'n_jobs': -1,
 'num_parallel_tree': 1,
 'reg_alpha': 0,
 'reg_lambda': 1,
 'scale_pos_weight': 1,
 'subsample': 1,
 'tree_method': 'exact',
 'validate_parameters': 1,
 'verbosity': 0)
pred_xgbf = classifier2.predict(X_val)
```

```
Out[ ]: 0.9475196230056412

accuracy = grid_search.best_score_
accuracy

0.9475196230056412
```

```
In [ ]: classifier2 = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                                colsample_bynode=1, colsample_bytree=0.5, gamma=0.2, gpu_id=-1,
                                importance_type='gain', interaction_constraints='',
                                learning_rate=0.2, max_delta_step=0, max_depth=15,
                                min_child_weight=1, missing=1, monotone_constraints=(),
                                n_estimators=100, n_jobs=-1, num_parallel_tree=1, random_state=0,
                                reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                                tree_method='exact', validate_parameters=1, verbosity=None)
classifier2.fit(X_train_new,y_train_new)
pred_xgbf = classifier2.predict(X_val)
```

### Validation Set

```
In [ ]: #Checking different metrics for decision tree model after tuning the hyperparameters
print("Checking different metrics for decision tree model after tuning the hyperparameters:\n")
print("Training accuracy: ",classifier2.score(X_val,y_val))
acc_score = accuracy_score(y_val, pred_xgbf)
print("Testing accuracy: ",acc_score)
conf_mat = confusion_matrix(y_val, pred_xgbf)
print("Confusion Matrix:\n",conf_mat)
roc_auc = roc_auc_score(y_val,pred_xgbf)
print("ROC AUC score: ",roc_auc)
class_rep2 = classification_report(y_val,pred_xgbf)
print("Classification Report:\n",class_rep2)
```

Checking different metrics for decision tree model after tuning the hyperparameters:

```
Training accuracy: 0.9365660262275084
Testing accuracy: 0.9365660262275084
Confusion Matrix:
[[2079 116]
 [ 92 952]]
ROC AUC score: 0.931140885440745
Classification Report:
              precision    recall  f1-score   support

    0         0.96         0.95         0.95         2195
    1         0.90         0.92         0.91         1084

 accuracy         0.93         0.93         0.94         3279
 macro avg         0.93         0.93         0.93         3279
 weighted avg         0.94         0.94         0.94         3279
```

### Test Set

```
In [ ]: pred_xgbf = classifier2.predict(X_test)

#Checking different metrics for decision tree model after tuning the hyperparameters
print("Checking different metrics for decision tree model after tuning the hyperparameters:\n")
print("Training accuracy: ",classifier2.score(X_test,y_test))
acc_score = accuracy_score(y_test, pred_xgbf)
print("Testing accuracy: ",acc_score)
conf_mat = confusion_matrix(y_test, pred_xgbf)
print("Confusion Matrix:\n",conf_mat)
roc_auc = roc_auc_score(y_test,pred_xgbf)
print("ROC AUC score: ",roc_auc)
class_rep2 = classification_report(y_test,pred_xgbf)
print("Classification Report:\n",class_rep2)
```

Checking different metrics for decision tree model after tuning the hyperparameters:

```
Training accuracy: 0.9350708733424783
Testing accuracy: 0.9350708733424783
Confusion Matrix:
[[1387  77]
 [ 65 658]]
ROC AUC score: 0.9287505951976055
Classification Report:
              precision    recall  f1-score   support

    0         0.96         0.95         0.95         1464
    1         0.90         0.91         0.90          723

 accuracy         0.93         0.93         0.94         2187
 macro avg         0.93         0.93         0.93         2187
 weighted avg         0.94         0.94         0.94         2187
```

## Feature Importance and Selection

### Feature Importance

```
In [ ]: fig, ax = plt.subplots(figsize=(18,10))
xgb.plot_importance(classifier2, max_num_features=50, height=0.5, ax=ax,importance_type='gain')
plt.show()
```



There are three methods to measure feature importances in xgboost.They are:

weight: The total number of times this feature was used to split the data across all trees. Cover: The number of times a feature is used to split the data across all trees weighted by the number of training data points that go through those splits. Gain: The average loss reduction gained when using this feature for splitting in trees. We used Gain in the above example and the model says when it used sentTransactions, the loss on average was reduced by 8.6%.

```
In [ ]: fig, ax = plt.subplots(figsize=(10,10))
xgb.plot_importance(classifier2, max_num_features=50, height=0.5, ax=ax,importance_type='weight')
plt.show()
```



When we considered weight,the model says that is used totalEtherBalance 1706 times to split the data across the trees.

### Feature Selection

```
In [ ]: model = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                            colsample_bynode=1, colsample_bytree=0.5, gamma=0.2, gpu_id=-1,
                            importance_type='gain', interaction_constraints='',
                            learning_rate=0.2, max_delta_step=0, max_depth=15,
                            min_child_weight=1, missing=1, monotone_constraints=(),
                            n_estimators=100, n_jobs=-1, num_parallel_tree=1, random_state=0,
                            reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                            tree_method='exact', validate_parameters=1, verbosity=None)
model.fit(X_train_new,y_train_new)
```

```
Out[ ]: * XGBClassifier
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=0.5,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, gamma=0.2, gpu_id=-1, grow_policy='depthwise',
              importance_type='gain', interaction_constraints='',
              learning_rate=0.2, max_bin=256, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=15, max_leaves=0, min_child_weight=1,
              missing=1, monotone_constraints=(), n_estimators=100, n_jobs=0,
              num_parallel_tree=1, predictor='auto', random_state=0,
              reg_alpha=0, reg_lambda=1, ...)
```

```
In [ ]: thresholds = sort(model.feature_importances_)
for thresh in thresholds:
    # select features using threshold
    selection = SelectFromModel(model, threshold=thresh, prefit=True)
    select_X_train = selection.transform(X_train_new)
    # train model
    selection_model = XGBClassifier()
    selection_model.fit(select_X_train, y_train_new)
    # eval model
    select_X_test = selection.transform(X_test)
    y_pred = selection_model.predict(select_X_test)
    predictions = (round(value) for value in y_pred)
    accuracy = accuracy_score(y_test, predictions)
    print("threshold=%.3f, n=%d, Accuracy: %.2f%%" % (thresh, select_X_train.shape[1], accuracy*100.0))
```

Threshold=0.000, n=18, Accuracy: 93.42%  
Threshold=0.032, n=17, Accuracy: 93.42%  
Threshold=0.034, n=16, Accuracy: 93.46%  
Threshold=0.035, n=15, Accuracy: 93.37%  
Threshold=0.035, n=14, Accuracy: 93.05%  
Threshold=0.035, n=13, Accuracy: 93.37%  
Threshold=0.036, n=12, Accuracy: 92.59%  
Threshold=0.045, n=11, Accuracy: 93.28%  
Threshold=0.047, n=10, Accuracy: 92.55%  
Threshold=0.051, n=9, Accuracy: 93.18%  
Threshold=0.052, n=8, Accuracy: 90.81%  
Threshold=0.056, n=7, Accuracy: 91.40%  
Threshold=0.059, n=6, Accuracy: 90.53%  
Threshold=0.069, n=5, Accuracy: 86.51%  
Threshold=0.087, n=4, Accuracy: 83.31%  
Threshold=0.095, n=3, Accuracy: 82.90%  
Threshold=0.112, n=2, Accuracy: 77.37%  
Threshold=0.118, n=1, Accuracy: 63.05%

From the results above if we want a balanced performance, accuracy tradeoff, we would go with only using n=11 with a threshold of 0.045 on the feature importances.

For feature importance we suggest at looking to combine multiple models with each other, or even using a model for each subsets of the datasets(one model for demographics, another model for transactions etc) Other complex models like anomaly detection neural networks can also be a good recommendation.

```
In [ ]:
```